

Non-Parametric Statistics Workshop 1

David Plazas Escudero
Juan Pablo Vidal
Juan Sebastián Cárdenas-Rodríguez
Mathematical Engineering, Universidad EAFIT

May 8, 2020

1 Workshop Exercises

The code done to solve this workshop can be found in a Jupyter notebook in [this link](#).

Exercise 1

The data used is the average daily temperatures in Canada for the past 35 years. The empirical cumulative distributions (ECDFs) for each year are presented in Figure 1, where the yellow-most curve represents the ECDF of the data recorded for the first year, and the blue-most is the last year.




Figure 1: Empirical distributions by years.

Based on this plots, it can be observed a climate change effect over the years. The plots of the first years arise for lower values of the temperature T , whereas the last years tend to be on the right of the plot. This can be interpreted as follows: since the plots for the early years are on the left, these years reported more data on lower average daily temperature than the most recent years. This means that recent years have been hotter in average daily temperature.

Exercise 2

The plug-in principle is a technique used in probability theory and statistics to estimate a parameter of a probability distribution (e.g., the expected value, the variance, a quantile) that cannot be

exactly computed. In general, the plug-in principle says that a feature of a given distribution can be approximated by the same feature of the empirical distribution of a sample of observations drawn from the given distribution [9].

The feature of the empirical distribution is called a plug-in estimate of the feature of the given distribution. For example, a quantile of a given distribution can be approximated by the analogous quantile of the empirical distribution of a sample of draws from the given distribution. The following is a formal definition of plug-in estimate.

A statistical functional $T(F)$ is any function of F . The plug-in estimator of $\theta = T(F)$ is defined by

$$\hat{\theta}_n = T(\hat{F}_n)$$

A functional of the form $\int a(x)dF(x)$ is called a linear functional. The plug-in estimator for linear functional $T(F) = \int a(x)dF(x)$ is:

$$T(\hat{F}_n) = \int a(x)d\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n a(X_i)$$

It is important to note that $T(F_n)$ converges to $T(F)$ as the sample size n increases.

In the practice, the limited area is calculated in the first quadrant for each empirical curve of the temperature data, this allows to extract the plug-in estimator of the mean, since the difference between the area with positive values and the area with negative values is extracted from the empirical curve, which allows estimating the mean of the series. The plug-in estimation of the mean for each year is presented in Figure 2, where a comparison with the natural estimator of the mean is presented.




Figure 2: Mean estimation of the series.

From Figure 2, an almost identical estimation can be seen between the two estimators, however it shows a slight difference between the mean estimates, this is because the only case in which both estimates are identical is when the sample size is infinite.




Figure 3: Comparison between the natural estimator and the plug-in estimator.

Exercise 3

Calculate and plot the confidence bands for the empirical continuous distribution function (ECDF) of the coldest and hottest year in average with a confidence of 95 %. Are there any sectors that are not enclosed in the bands?

Proof. Let n be the size of the sample and $1 - \alpha$ the desired confidence for the bands. This method is based on the Dvoretzky–Kiefer–Wolfowitz inequality (see [10]). In order to calculate the confidence bands for the ECDF, we first define ϵ_n by the following formula:

$$\epsilon_n = \sqrt{\frac{1}{2n} \ln \left(\frac{2}{\alpha} \right)}$$

Let $\hat{F}_n(x)$ be the ECDF. Then, for each x in the ECDF we define the lower ($L(\cdot)$) and upper ($U(\cdot)$) bound by:

$$L(x) = \max\{\hat{F}_n(x) - \epsilon_n, 0\}$$

$$U(x) = \min\{\hat{F}_n(x) + \epsilon_n, 1\}$$

The results obtained by using the temperatures of the coldest and hottest year are seen in Figure 4.




Figure 4: Bands for the coldest and hottest year.

It can be seen that the upper and lower bands fully enclose the ECDF. Nevertheless, there exists two points where the function and the bands meets. This happens in the lower band at the beginning and the upper band at the final point.

This phenomenon is due to the full certainty at those points in a sense that the lower bound, at the start, has to be the same point as it cannot go lower than 0. A similar reasoning can explain the upper bound and the final point. \square

Exercise 4

Write and execute a code that allows to visualize the Glivenko Cantelli for a Weibull distribution.

Proof. The Glivenko Cantelli shows the relationship of how the difference between the empirical and theoretical distribution changes depending of the number of the sample.

In this manner, if n is the size of the sample, $\hat{F}_n(\cdot)$ is the ECDF of that sample and $F(\cdot)$ is the theoretical distribution the theorem states that:

$$\sup_x |\hat{F}_n(x) - F(x)| \xrightarrow{\text{a.s.}} 0.$$

We generated Weibull random variables of different sizes n_i calculated with:

$$n_i = 2^i, \quad \text{for } i = 1, \dots, 20$$

The results of the experiment can be found in Figure 5.




Figure 5: Glivenko Cantelli theorem visualization.

It is clear that the biggest difference between the ECDF and the theoretical distribution starts to go to 0 when n gets bigger. Furthermore, although the behavior in the graph does not show a monotone decrease, it does display the asymptotic behavior described in the theorem. \square

Exercise 5

Theorem 1.1 (Jensen's Inequality). Let $g(\cdot)$ be a concave function and let X be a random variable, then

$$\mathbb{E}[g(X)] \leq g(\mathbb{E}[X]).$$

Proof. Let $g(\cdot)$ be a concave function and $L(x)$ be its tangent line at a fixed point x_0 . It holds that (see Theorem 6 from [6])

$$g(x) \leq L(x). \quad (1)$$

If $x_0 = \mathbb{E}[X]$. From (1) and given that $\mathbb{E}[\cdot]$ preserves order, we have

$$\mathbb{E}[g(X)] \leq \mathbb{E}[L(X)] = \mathbb{E}[a + bX] = a + b\mathbb{E}[X] = L(\mathbb{E}[X]) = g(\mathbb{E}[X]).$$

□

1.1 Exercise 6

L-estimator

An L-estimator is a linear combination of order statistics of the measurements, often called L-statistic. They have the advantage of being robust and simple to use, but tend to have problems with low efficiency. Not all L-estimators are robust (for example, the minimum, maximum and mean are not considered robust) and some have better efficiency than others. L-statistics with multiple (correct) weights will tend to be more efficient than those with fewer weights or poorly chosen weights [1].

L-estimators for a sample size n are defined by T_n :


$$T_n(X_1, \dots, X_n) = \sum_{i=1}^n a_i X_{n,(i)}$$

$$a_i = \int_{(i-1)/n}^{i/n} h(x) dx, \quad \int_0^1 h(x) dx = 1$$


where: X_n are the order statistics and a_i are weight factors.

L-estimators are useful in robust statistics, but they are considered inefficient. In modern statistics M-estimators are preferred, though these are much more difficult to implement computationally. In many circumstances, L-estimators are reasonably easier to implement and, thus, adequate for initial estimation. M-estimators provide robust statistics that also have high relative efficiency, at the cost of being much more computationally complex [1].

Examples



(a) Median estimator.



(b) Mean estimator.

Figure 6: L-estimators.

Figure 6 shows the influence function of the median and mean. It can be seen that the median is a much more robust statistic, which maintains a high breakdown point, above 50%, allowing the estimator to remain unchanged when modifying or contaminating the sample. In this case, random data is generated from a standard normal distribution, and by contaminating the sample in multiple percentages, the median remains with a very similar value until it exceeds the breakdown point, which is 50%, the highest in the L-estimators.

On the other hand, it can be seen that the mean, a measures of central tendency, is a non-robust estimator, which varies according to how the sample is modified, with a breakdown point of $1/N$, since any change modifies its value.




Figure 7: L-estimators.

Figure 7 shows the influence function of the maximum and minimum, these non-robust estimators have a break down point of 0, since any modification of the sample can alter its value considerably, as can be seen in both figures, where in any percentage of contamination, the estimator has a totally different value.

M-estimator

M-estimators are a broad class of extremum estimators for which the objective function is a sample average [1]. The extremum estimators are a wide class of estimators for parametric models that are calculated through maximization (or minimization) of a certain objective function. An M-estimator minimizes the function:

$$Q(e_i, \rho) = \sum_i \rho\left(\frac{e_i}{s}\right)$$


where ρ is a symmetric function of the residuals

- The effect of ρ is to reduce the influence of outliers
- s is an estimate of scale.
- The robust estimates $\hat{\beta}$ are computed by the iteratively re-weighted least squares algorithm
- We have several choices available for the weighting functions to be used


The M-estimator is more efficient under certain conditions: the data contains y outliers and the model matrix is measured with no errors [8].

This estimators are especially useful when the data has outliers or is contaminated by heavy tailed errors. On the other hand, M-estimation is not recommended when anomalous data reflects the true population, or the population is made up of distinct mixture of distributions.

Examples




(a) Andrew's Wave estimator.




(b) Tukey's Biweight estimator.

Figure 8: M-estimators.



(a) Hampel's estimator.



(b) Huber's T estimator.

Figure 9: M-estimators.

Figures 8 and Figure 9 shows redescending M-estimators, which are estimators which have functions that are non-decreasing near the origin, but decreasing toward 0 far from the origin. These kinds of estimators are very efficient, have a high breakdown point (close to 0.5) and, unlike other outlier rejection techniques, they do not suffer from a masking effect. They are efficient since they completely reject gross outliers, and do not completely ignore moderately large outliers (like median). In this case the worst estimator is Hubert T estimator, because the other estimators completely reject gross outliers, while the Huber estimator effectively treats these the same as moderate outliers [8].

Exercise 7

Deduce the distribution and density of the j -th ordered statistic. Explain with detail what would be a simple procedure to simulate the j -th ordered statistic. Simulate 1000 observations of some ordered statistic of a sample of size n that comes from a Weibull distribution. Draw in a same graph the ECDF and theoretical distribution.

Proof. Let's first deduce the distribution and density of the j -th ($X_{[j]}$) ordered statistic. Let X_1, \dots, X_n be independent random variables that come from a distribution $F(\cdot)$. Hence, the probability that $X_i \leq t$ is given by:

$$P(X_i \leq t) = F(t)$$

Let Z_t be a random variable that represents the number of variables whose value are less than t . Hence:

$$Z_t \in \{0, 1, \dots, n\}$$

Hence:

$$\begin{aligned} P(Z_t = 0) &= P(X_1 > t \wedge \dots \wedge X_n > t) \\ &= P(X_1 > t) \dots P(X_n > t) \\ &= (1 - F(t)) \dots (1 - F(t)) \\ &= (1 - F(t))^n \\ P(Z_t = 1) &= \binom{n}{1} P(X_1 \leq t \wedge X_2 > t \wedge \dots \wedge X_n > t) \\ &= \binom{n}{1} F(t)(1 - F(t)) \dots (1 - F(t)) \\ &= \binom{n}{1} F(t)(1 - F(t))^{n-1} \\ &\vdots \\ P(Z_t = j) &= \binom{n}{j} F(t)^j (1 - F(t))^{n-j} \end{aligned}$$

Hence, the distribution j -th ordered statistic is given by:

$$P(X_{[j]} \leq t) = P(Z_t \geq j) = \sum_{i=j}^n \binom{n}{i} F(t)^i (1 - F(t))^{n-i}$$

In these manner, we obtained the distribution for the j -th ordered statistic. Then, to obtain the density ($f_{[j]}(\cdot)$) we just differentiate hence:

$$\begin{aligned} f_{[j]}(t) &= \frac{d}{dt} P(X_{[j]} \leq t) \\ &= \sum_{i=j}^n \binom{n}{i} [iF(t)^{i-1}(1 - F(t))^{n-i}f(t) - (n-i)F(t)^i(1 - F(t))^{n-i-1}f(t)] \\ &= f(t) \sum_{i=j}^n \binom{n}{i} F(t)^{i-1}(1 - F(t))^{n-i-1} [i(1 - F(t)) - (n-i)F(t)] \\ &= f(t) \sum_{i=j}^n \binom{n}{i} F(t)^{i-1}(1 - F(t))^{n-i-1} (i - nF(t)) \\ &= f(t) \left[\sum_{i=j}^n i \binom{n}{i} F(t)^{i-1}(1 - F(t))^{n-i-1} - \sum_{i=j}^n n \binom{n}{i} F(t)^i(1 - F(t))^{n-i-1} \right] \end{aligned}$$

$$\begin{aligned}
&= n f(t) \left[\sum_{i=j}^n \binom{n-1}{i-1} F(t)^{i-1} (1 - F(t))^{n-i-1} - \sum_{i=j}^n \binom{n}{i} F(t)^i (1 - F(t))^{n-i-1} \right] \\
&= n \binom{n-1}{j-1} F(t)^{j-1} (1 - F(t))^{n-j} f(t)
\end{aligned}$$

Furthermore, simulating a j -th statistic can be done easily. In first place, generate samples of size n and obtain the j -th statistic of that sample. These generates one data of the j -th ordered statistic. Repeat the previous process until the desired number of data is obtained.

The result of the simulation of the 5-th ordered statistic for a Weibull distribution can be found in Figure 10.




Figure 10: Comparison between ECDF and theoretic distribution.

□

Exercise 8

Suppose X is an exponentially random variable of parameter β . Calculate:

$$P(|X - \mu| > k\sigma)$$

for $k > 1$.

Proof. Recall Chebyshev's inequality: let Y be a random variable and let $\mathbb{E}[Y] = \mu$ and $\text{Var}[Y] = \sigma^2$, then

$$P(|Y - \mu| \geq t) \leq \frac{\sigma^2}{t^2}.$$

Let

$$\begin{aligned}
P(|X - \mu| > k\sigma) &= P(-k\sigma > X - \mu > k\sigma) \\
&= P(-k\sigma + \mu > X > k\sigma + \mu) \\
&= P(X < \mu - k\sigma) + P(X > k\sigma + \mu) \\
&= P\left(X < \frac{1-k}{\beta}\right) + 1 - P\left(X < \frac{k+1}{\beta}\right) \quad \text{but } 1-k < 0
\end{aligned}$$

$$\begin{aligned}
&= 1 - \left(1 - e^{-\frac{k+1}{\beta^2}}\right) \\
&= P(|X - \mu| > k\sigma) \\
&\leq \frac{\sigma^2}{(k\sigma)^2} \\
&\leq \frac{1}{k^2}
\end{aligned}$$

Clearly $e^{-\frac{k+1}{\beta^2}}$ is always lesser than 1 and because $k > 1$ then $\frac{1}{k^2}$ is also lesser than one. Thus, $P(|X - \mu| > k\sigma)$ is bounded by Chebyshev's inequality. \square

Exercise 9

Prove that if $X \sim \text{Poisson}(\lambda)$, then

$$P(X \geq 2\lambda) \leq \frac{1}{\lambda}.$$

Proof. Let $X \sim \text{Poisson}(\lambda)$. Recall Chebyshev's inequality: let Y be a random variable and let $\mathbb{E}[Y] = \mu$ and $\text{Var}[Y] = \sigma^2$, then

$$P(|Y - \mu| \geq t) \leq \frac{\sigma^2}{t^2}.$$

Clearly, $\mathbb{E}[X] = \text{Var}[X] = \lambda$. If we set $t = \lambda$, then

$$\begin{aligned}
P(|X - \lambda| \geq \lambda) &\leq \frac{1}{\lambda} \\
P[(X - \lambda \geq \lambda) \cup (X - \lambda \leq -\lambda)] &\leq \frac{1}{\lambda} \\
P[(X \geq 2\lambda) \cup (X \leq 0)] &\leq \frac{1}{\lambda} \\
P(X \geq 2\lambda) + P(X \leq 0) &\leq \frac{1}{\lambda} \\
P(X \geq 2\lambda) &\leq \frac{1}{\lambda}
\end{aligned}$$

\square

Exercise 10

Definition 1.1. The sequence $\{X_n\}$ of random variables is said to converge in probability to X if

$$\forall \epsilon > 0, \lim_{n \rightarrow \infty} P(|X_n - X| > \epsilon) = 0$$

Definition 1.2. The sequence $\{X_n\}$ of random variables is said to converge in mean-square to X if

$$\lim_{n \rightarrow \infty} \mathbb{E}[(X_n - X)^2] = 0$$

Theorem 1.2. The convergence in mean-square implies convergence in probability.

Proof. Let $\{X_n\}$ be a sequence of random variables that converges in mean-square to X . Recall Markov's inequality: Let Y be a non-negative random variable and suppose $\mathbb{E}[Y]$ exists. Then for any $t > 0$,

$$P(Y > t) \leq \frac{\mathbb{E}[Y]}{t}.$$

Take $Y = |X_n - X|$ and $t = \epsilon$, then

$$\begin{aligned} P(|X_n - X| > \epsilon) &= P[(X_n - X)^2 > \epsilon^2] \\ &\leq \frac{\mathbb{E}[(X_n - X)^2]}{\epsilon^2} \end{aligned}$$

Since $\{X_n\}$ converges in mean-square to X , $\mathbb{E}[(X_n - X)^2] \rightarrow 0$ as $n \rightarrow \infty$, which directly implies that $P(|X_n - X| > \epsilon) \rightarrow 0$ as $n \rightarrow \infty$. \square

Exercise 11

Show that the ECDF converges in probability to the theoretical continuous distribution function.

Proof. Let X_i , for $i = 1, \dots, n$, be a independent data sample. Then, the empirical distribution function is defined as

$$\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n I(X_i \leq x)$$

where

$$I(X_i \leq x) = \begin{cases} 1 & \text{if } X_i \leq x \\ 0 & \text{otherwise} \end{cases}$$

Hence, to see that it converges in probability lets see if the MSE tends to 0 when the number of samples is bigger. Let's calculate the expectancy of the estimator.

$$\begin{aligned} \mathbb{E}[\hat{F}_n(x)] &= \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n I(X_i \leq x)\right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[I(X_i \leq x)] \\ &= \frac{1}{n} \sum_{i=1}^n P(X_i \leq x) \\ &= \frac{1}{n} \sum_{i=1}^n F(x) \\ &= F(x) \end{aligned}$$

In this manner, it is a non-biased estimator for the theoretical distribution. Hence, the MSE is calculated by the variance. Then:

$$\text{MSE} = \text{Var}[\hat{F}_n(x)] + \text{Bias}(\hat{F}_n(x), F(x))$$

$$\begin{aligned}
&= \text{Var} \left[\frac{1}{n} \sum_{i=1}^n I(X_i \leq x) \right] \\
&= \frac{1}{n^2} \sum_{i=1}^n \text{Var} [I(X_i \leq x)] \\
&= \frac{1}{n^2} \sum_{i=1}^n [\text{P}(X_i \leq x)(1 - \text{P}(X_i \leq x))] \\
&= \frac{1}{n^2} \sum_{i=1}^n [F(x)(1 - F(x))] \\
&= \frac{F(x)(1 - F(x))}{n}
\end{aligned}$$

Hence, when $n \rightarrow \infty$ the MSE tends to 0. Therefore,

$$\hat{F}_n(x) \xrightarrow{P} F(x)$$

□

Exercise 12

Consider the daily temperatures of the hottest year in average. Calculate a confidence interval for the maximum temperature. Calculate the bias of $T_{[n]}$ and the variance.

Proof. To calculate the interval for the maximum temperature, we applied the normal bootstrap method. The normal bootstrap method consists of, in a sample X_1, \dots, X_n and a confidence of $1 - \alpha$:

1. First, extract k samples of the size n with repetition and calculate the desired statistic to find its confidence intervals. Let S_i be the sample of each statistic calculated.

2. Calculate:

$$v_{\text{boot}} = \frac{1}{k} \sum_{i=1}^k \left(S_i - \frac{1}{n} \sum_{i=1}^k S_i \right)^2$$

3. Then, the interval of confidence for the statistic is, with S the statistic calculated in the original sample:

$$(S - z_{\alpha/2} \sqrt{v_{\text{boot}}}, S + z_{\alpha/2} \sqrt{v_{\text{boot}}})$$

where $z_{\alpha/2}$ is the $1 - \alpha/2$ quantile of a standard normal distribution.

The result for the maximum temperatures interval is, with $\alpha = 0.05$ and $k = 1000$:

$$T_{[n]} \in (18.486, 18.714)$$

On the other hand, to calculate the variance and bias the Jackknife method was used. The Jackknife method consists is explained in the following exercise. Hence, we find that the bias and the variance and we obtain:

$$\begin{aligned}
b_{\text{jack}} &= -0.10 \\
v_{\text{jack}} &= 7.49 \cdot 10^{-8}
\end{aligned}$$

□

Exercise 13

In this section, a sample of a uniform distribution in the interval $[0,1]$ is generated. Then a bootstrap method to calculate the variance and a Jackknife method to calculate the bias of this sample are implemented.

The bootstrap and the jackknife are non-parametric methods for computing standard errors and confidence intervals [7]. A short review of the methodology will be presented, based on [10].

The Jackknife

Jackknife is a simple method for approximating the bias and variance of an estimator. Let $T_n = T(X_1, \dots, X_n)$ be an estimator of some quantity θ and let $\text{bias}(T_n) = \mathbb{E}(T_n) - \theta$ denote the bias. Let $T_{(-i)}$ denote the statistic with the i^{th} observation removed. The jackknife bias estimate is defined by

$$b_{\text{jack}} = (n-1)(\bar{T}_n - T_n)$$

where $\bar{T}_n = n^{-1} \sum_i T_{(-i)}$. The bias-corrected estimator is $T_{\text{jack}} = T_n - b_{\text{jack}}$

The Bootstrap

Bootstrap is a method for estimating the variance and the distribution of a statistic $T_n = g(X_1, \dots, X_n)$. Let $\mathbb{V}_F(T_n)$ denote the variance of T_n . We have added the subscript F to emphasize that the variance is a function of F . If we knew F we could, at least in principle, compute the variance. For example, if $T_n = n^{-1} \sum_{i=1}^n X_i$ then

$$\mathbb{V}_F(T_n) = \frac{\sigma^2}{n} = \frac{\int x^2 dF(x) - (\int x dF(x))^2}{n}$$

which is clearly a function of F .

With the bootstrap, we estimate $\mathbb{V}_F(T_n)$ with $\mathbb{V}_{\hat{F}_n}(T_n)$. In other words, we use a plug-in estimator of the variance. since, $\mathbb{V}_{\hat{F}_n}(T_n)$ may be difficult to compute, we approximate it with a simulation estimate denoted by v_{boot}

Its important to know that The jackknife is a linear approximation of the bootstrap.

In the simulation, the Variance Bootstrap obtained from the uniform sample of the Order statistic is $v_{\text{boot}} = 3.5815 \times 10^{-8}$, this indicates that almost all of the Order data values are nearly identical. Moreover, the theoretical bias and the Jackknife bias generated similar results, $b_{\text{jack}} = -1.6166 \times 10^{-5}$ and $b_{\text{theoretical}} = -9.9990 \times 10^{-5}$, which shows that the bias of the estimator is unbiased, hence the sampling distribution has a mean that is equal to the parameter being estimated.

It should be noted that the jackknife method is less computationally expensive, but the bootstrap has some statistical advantages, such as its simplicity to derive estimates of standard errors and confidence intervals for complex estimators of complex parameters of the distribution.

Exercise 14

Show the differences between the parametric and non-parametric bootstrap. Investigate robust versions of the bootstrap method and show examples about their performance.

Proof. Following the ideas from [10] and [5], Bootstrap is performed over a sample X_1, \dots, X_n . In the case of the nonparametric Bootstrap, we generate new samples (resampling) based on the ECDF \hat{F}_n of the sample; the generation of samples from the ECDF is equivalent to draw samples X_1^j, \dots, X_n^j from the original data **with replacement** for $j = 1, \dots, N$, where N is the number of Bootstrap resamples. On the other hand, the parametric Bootstrap takes into consideration that the data comes from an specific distribution F_θ that, clearly, depends on an unknown parameter θ ; instead of drawing from \hat{F}_n , we draw from $F_{\hat{\theta}}$, where $\hat{\theta}$ is an estimator of θ based on the sample. This method is just as accurates as the nonparametric, but under certain scenarios could not behave properly. An excellent example of the parametric and nonparametric Bootstrap is presented in point 11 of page 40 from [10]. □

Exercise 15

The data used in this section is bivariate, taking the temperature from the years that are, in average, the coldest and hottest. Let $\mathbf{Y} \in \mathbb{R}^{n \times 2}$ be the set of bivariate samples of size n (data matrix). In our case, $n = 365$ days.

The usual cutout of outliers in elliptical data (e.g. bivariate normally distributed data) is made using the χ^2 distribution. However, the bivariate data obtained from the temperatures is not guaranteed to follow such elliptical behavior. Therefore, the usual cutout will be done with the same methodology as the robust variations. The methods for estimating the covariance matrix will be first presented, then the methodology to remove outliers will be described and finally the results will be presented.

Usual Estimation

Definition 1.3. Let \mathbf{x} be an multivariate observation from a set of observations with mean vector $\boldsymbol{\mu}$ and covariance matrix Σ . The Mahalanobis distance of the observation \mathbf{x} is

$$d^2(\mathbf{x}) = (\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}). \quad (2)$$

Clearly, the usual cutout is performed using the natural estimation of the Mahalanobis distance, using the standard unbiased estimators of the mean vector and covariance matrix from the data matrix \mathbf{Y} , presented respectively in Equation (3), where \mathbf{y}_i is the i -th row of the data matrix \mathbf{Y} .

$$\hat{\boldsymbol{\mu}} = \bar{\mathbf{y}} = \left[\frac{1}{n} \sum_{i=1}^n \mathbf{y}_i \right]^T, \quad \hat{\Sigma} = S = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{y}_i^T - \bar{\mathbf{y}})(\mathbf{y}_i^T - \bar{\mathbf{y}})^T \quad (3)$$

Comedian Estimation

The first robust calculation of the Mahalanobis distance is based on a robust estimation of the covariance matrix following the ideas from [3], using the following definition.

Definition 1.4. Let X and Y be random variables. The comedian between X and Y is defined as

$$\text{Com}(X, Y) = \text{Med} [(X - \text{Med}(X))(Y - \text{Med}(Y))]$$

The covariance matrix is then estimated by applying the comedian to each entry. Then the Mahalanobis distance formula is applied using this matrix and the median vector instead of the mean.

Kendall Estimation

This method uses the Kendall rank correlation coefficient, usually known as Kendall's τ coefficient, originally proposed in [4]. Each entry of the covariance matrix is estimated using

$$\text{Cov}(X, Y) = \rho_k S_X S_Y,$$

where ρ_k is Kendall's τ coefficient, and S_X and S_Y are the respective standard deviations.

Spearman Estimation

This estimation was performed similarly to Kendall's. The covariance matrix is estimated using

$$\text{Cov}(X, Y) = \rho_s S_X S_Y,$$

where ρ_s is Spearman's correlation coefficient (see [2]), and S_X and S_Y are the respective standard deviations.

Cutout Procedure

As previously mentioned, the cutout method is performed with the same procedure for all the covariance matrix estimations. The procedure is presented as follows:

1. Estimate the covariance matrix accordingly to the estimation method.
2. Apply equation (2) to each bivariate data, with the estimation of the covariance matrix from step 1 and the corresponding estimated mean vector.
3. Fit a continuous distribution to the vector of squared distances: take the available continuous distributions from Python's package SciPy, fit each distribution to the squared distances and apply a Kolmogorov-Smirnov goodness-of-fit test in order to rank the fitted distributions. Keep the best fit.
4. Calculate the $\alpha/2$ and $1 - \alpha/2$ quantiles of the fitted distribution.
5. Mark the outliers as the data points that are beyond the quantiles obtained on step 4.

Results

In Figure 11 shows the scatter plot of the bivariate data used in this section.




Figure 11: Scatter plot of bivariate data.

The first cut method was done with the usual estimation of the covariance matrix. In Figure 12, the results for this procedure are presented. The best-fit was a Mielke's Beta-Kappa distribution, whose probability density function(p.d.f.) is

$$f(t) = \frac{kt^{k-1}}{(1+t^s)^{1+\frac{k}{s}}}, \quad t > 0$$

with parameters $k = 0.96$, $s = 3.85$, $\text{loc}=0.01$ and $\text{scale}=3.37$, and it was not rejected with a p-value of 0.34.




Figure 12: Results using the usual estimation.

On the other hand, the first robust approach is using the Comedian procedure above described, where each entry is estimated using the comedian and the mean vector is replaced with the median vector. The results are shown in Figure 13. The best-fit was an inverse gaussian distribution, whose density is given by

$$f(t) = \frac{1}{\sqrt{2\pi t^3}} e^{-\frac{(t-\mu)^2}{2t\mu^2}}, \quad t > 0$$

with parameters $\mu = 1.11$, $\text{loc}=-254.79$ and $\text{scale}=2667.77$, and it not rejected with a p-value of 0.28.




Figure 13: Results using the comedian estimation.

Furthermore, the same procedure was carried out using the covariance matrix estimation based on Kendall's τ correlation coefficient. Figure 14 shows the obtained results for the bivariate data. The

best-fit was a folded normal distribution, whose p.d.f is given by

$$f(t) = \sqrt{\frac{2}{\pi}} \cosh(ct) e^{-\frac{(t^2+c^2)}{2}}, \quad t > 0$$

with parameters $\mu = 1.01$, $\text{loc}=0$ and $\text{scale}=2.19$, and it was not rejected with a p-value of 0.45.




Figure 14: Results using Kendall estimation.

Finally, the last robust estimation of the covariance matrix was constructed using Spearman's correlation coefficient and the outlier detection results are presented in Figure 15. The best-fit was a Gompertz distribution, whose p.d.f. is given by

$$f(t) = ce^t e^{-ce^t - 1}, \quad t > 0$$

with parameters $c = 1.04$, $\text{loc}=0$ and $\text{scale}=2.19$, and it was not rejected with a p-value of 0.31.






Figure 15: Results using Spearman estimation.

In the following results, the same procedures of outlier cutout is presented, but taking a "damaged" sample. This new sample takes 30 random bivariate observations and adds normally distributed observations, with mean 10 and standard deviation of 0.5. The results for the usual, the comedian, the Kendall and the Spearman cutouts are presented, respectively, in Figures 16, 17, 18 and 19. It is important to highlight that the fitting was done assuming the same distributions from the previous procedures.




(a) Fitted distribution.




(b) Outliers cutout.

Figure 16: Results using the usual estimation on contaminated data.




(a) Fitted distribution.




(b) Outliers cutout.

Figure 17: Results using the comedian estimation on contaminated data.



(a) Fitted distribution.



(b) Outliers cutout.

Figure 18: Results using Kendall estimation on contaminated data.




Figure 19: Results using Spearman estimation on contaminated data.

2 Book Exercises

All exercises in this section are extracted from [10].

2.1 Page 10, Exercise 2

Let $X_1, \dots, X_n \sim \text{Bernoulli}(p)$. A pointwise asymptotic $1 - \alpha$ confidence interval for p is

$$\hat{p}_n \pm z_{\alpha/2} \sqrt{\frac{\hat{p}_n (1 - \hat{p}_n)}{n}}$$

where $\hat{p}_n = n^{-1} \sum_{i=1}^n X_i$. It follows from Hoeffding's inequality that a finite sample confidence interval is

$$\hat{p}_n \pm \sqrt{\frac{1}{2n} \log \left(\frac{2}{\alpha} \right)}$$

In Figure 20 is presented the coverage and length of the pointwise asymptotic $1 - \alpha$ confidence interval and the finite sample confidence interval (Hoeffding) to estimate the p parameter of a Bernoulli distribution. To simulated this scenario a Bernoulli distribution was generate with parameters $p = 0.2$ and $\alpha = .05$.




Figure 20: Confidence bands for the value of p .


The result shows that an large sample must be generated so that the pointwise interval has accurate coverage of the p , it can also be seen that both intervals have similar length when reaching an

accurate coverage, being Hoeffding the interval with the highest coverage due to a very slight difference.


Note that the sample sizes shown on the graph are in powers.

Page 24, Exercise 3

This exercise uses a sample of 100 observations that are normally (standard) distributed. It is required to compute a 95% confidence band for the CDF F and how many times this band covers the true CDF of the data, by repeating the process 1000 times. For this procedure we use the bands derived from Dvoretzky–Kiefer–Wolfowitz inequality [10]. In Figures 21a and 21b we present the worst and best band obtained over the experiments, respectively. The experiments yield that only 8 out of the 1000 bands simulated contain the true CDF of the standard normal random variable.



(a) Worst enclosing band.




(b) Best enclosing band.

Figure 21: Confidence bands for F of normal data.


Additionally, the experiment was repeated with a Cauchy distribution with p.d.f. given by

$$f(t) = \frac{1}{\pi(1+t^2)}, \quad \forall t \in \mathbb{R}.$$

The obtained result are presented in Figure 22, as well with the worst and best enclosing bands.



(a) Worst enclosing band.



(b) Best enclosing band.

Figure 22: Confidence bands for F of normal data.

Page 24, Exercise 7

Let $X_1, \dots, X_n \sim \text{Bernoulli}(p)$ and let $Y_1, \dots, Y_m \sim \text{Bernoulli}(q)$ with a $p = 0.5$ and $q = 1 - p$.

The plug-in estimator for p is $\hat{p} = 0.51$ with a estimated standard error $\hat{se} = 0.0928$ and an approximate 90 percent confidence interval for p of $(0.3405, 0.6594)$.

The plug-in estimator for $p - q$ is $\widehat{p - q} = 0.5$ with a estimated standard error $\hat{se} = 0.0832$ and an approximate 90 percent confidence interval for $p - q$ of $(0.3830, 0.6169)$.

The estimates above show very similar results, with very slight differences in the intervals and the estimated standard error.

Page 39, Exercise 2

The influence function is used to approximate the standard error of a plug-in estimator. In this case, the empirical influence function was use:

The empirical influence function is defined by $\hat{L}(x) = L_{\hat{F}_n}(x)$. Thus,

$$\hat{L}(x) = \lim_{\epsilon \rightarrow 0} \frac{T\left((1 - \epsilon)\hat{F}_n + \epsilon\delta_x\right) - T\left(\hat{F}_n\right)}{\epsilon}$$

To estimate the the standard error, the empirical influence function, the jackknife and the bootstrap were used. The results were: Variance Bootstrap $v_{\text{boot}} = 7.5380 \times 10^{-17}$, Variance Jackknife $v_{\text{jack}} = 0.0176$, Variance Influence $v_{\text{Influence}} = 0.4731$ and a 95 percent studentized pivotal bootstrap confidence interval of $[0.5459189161795877, 0.5459189161795885]$. This shows that the LSAT scores and GPA data are very similar to each other, since a very close variance is observed in the multiple test cases.

Page 40, Exercise 6

The objective of this experiment is to compare the four different confidence intervals exposed for the Bootstrap method. The Bootstrap is performed over a sample of lognormally distributed data, taking 50 standard normal observations and making the exponentiation with the natural base. The statistic to be evaluated is the skewness of the data, which is well known to be positive for lognormal data. The results are presented in Table 1.

Method	95% C.I.
Normal	(1.23, 3.41)
Pivotal	(1.07, 3.22)
Studentized	(1.63, 3.57)
Percentile	(1.05, 3.20)

Table 1

Page 40, Exercise 7

Let

$$X_1, \dots, X_n \sim t_3$$

where $n = 25$ and t_3 is a student's t-distribution with $\nu = 3$. Let $\theta = T(F) = (q.75 - q.25)/1.34$. A simulation was implemented to compare the coverage and length of the following confidence intervals for θ : Normal interval with standard error from the bootstrap : $[1.5486, 0.1564]$, bootstrap percentile interval: $[0.3823, 1.3026]$ and a normal interval with standard error from the jackknife: $[0.8141, 0.8905]$.

This results shows that the jackknife does not give a consistent estimator of the variance of a quantile, showing a very short coverage for the estimator, however, both bootstrap methods show a coverage and length very similar to each other, with the difference that the bootstrap percentile interval has a smaller lower bound, this is because it is a quantile method, With a statistics that have a small lower bound.

Page 40, Exercise 10

Let $X_1, \dots, X_n \sim \text{Normal}(\mu, 1)$. Let $\theta = e^\mu$ and $\hat{\theta} = e^{\bar{x}}$ be the MLE. Create a data set (using $\mu = 5$) consisting of $n = 100$ observations.

- Use the delta method to get the se and 95 percent confidence interval for θ . Use the parametric bootstrap to get the se and 95 percent confidence interval for θ . Use the nonparametric bootstrap to get the se and 95 percent confidence interval for θ . Compare your answers.
- Plot a histogram of the bootstrap replications for the parametric and nonparametric bootstraps. These are estimates of the distribution of $\hat{\theta}$. The delta method also gives an approximation to this distribution, namely $\text{Normal}(\hat{\theta}, \hat{se}^2)$. Compare these to the true sampling distribution of $\hat{\theta}$. Which approximation is closer to the true distribution?

Proof. Let's solve each item.

- For the delta method, it is necessary to find the influence function for the estimator. In first place, the plug-in estimator:

$$\begin{aligned}\theta &= e^\mu \\ &= T(F) \\ &= e^{\int u dF(u)}\end{aligned}$$

Now, the influence function:

$$\begin{aligned}L_F(x) &= \lim_{\epsilon \rightarrow 0} \frac{T[(1-\epsilon)F + \epsilon\delta_x] - T(F)}{\epsilon} \\ &= \lim_{\epsilon \rightarrow 0} \frac{e^{\int u d[(1-\epsilon)F + \epsilon\delta_x]} - e^{\int u dF(u)}}{\epsilon} \\ &= \lim_{\epsilon \rightarrow 0} \frac{e^{(1-\epsilon)\mu + \epsilon x} - e^\mu}{\epsilon} \\ &= e^\mu \lim_{\epsilon \rightarrow 0} \frac{e^{\epsilon(x-\mu)} - 1}{\epsilon} \quad \begin{bmatrix} 0 \\ 0 \end{bmatrix}\end{aligned}$$

Applying L'Hopital's rule

$$= e^\mu \lim_{\epsilon \rightarrow 0} \left((x - \mu) e^{\epsilon(x-\mu)} \right)$$

$$= e^{\mu}(x - \mu)$$

Hence:

$$\hat{L}(x) = e^{\bar{X}}(x - \mu)$$

Using the parametric and non-parametric normal bootstrap method, the results are found in Table 2.

Method	95% C.I.
Delta	(128.52, 188.11)
Parametric	(127.60, 189.03)
Non-Parametric	(128.25, 188.38)

Table 2

- b) The comparison of the real distribution and the estimation done in the methods is shown in Figure 23




Figure 23: Comparison of distributions.

It is easy to see that the non-parametric bootstrap method is the closest one to the real distribution, as its peak does not have a significant bias, different to the other ones.

□

Page 41, Exercise 11

Let $X_1, \dots, X_n \sim \text{Uniform}(0, 1)$, the maximum likelihood estimator (MLE) for $\theta = 1$ is $\hat{\theta} = X_{\max} = \max\{X_1, \dots, X_n\}$. The distribution of $\hat{\theta}$ is a direct consequence of the proof made in Subsection 1.1 and it is given by

$$F(t) = t^n.$$

In Figure 24 shows the theoretical density of X_{\max} and the approximation in histograms by parametric and nonparametric Bootstrap.




Figure 24: Density of X_{\max} , parametric and nonparametric histograms of X_{\max} .

Let $\hat{\theta}^*$ be the estimation of the parameter θ for each Bootstrap replication. For a parametric Bootstrap, one could draw samples from $F_{\hat{\theta}}$, instead of drawing samples from the empirical distribution \hat{F}_n . Therefore, we would draw a N (number of Bootstrap iterations) new samples $X_1^j, \dots, X_n^j \sim \text{Uniform}(0, X_{\max})$, for $j = 1, \dots, N$. Clearly, $P(\hat{\theta}^* = \hat{\theta}) = 0$, since it is the probability of a point on a continuous random variable, which is obvious that it has a probability measure 0.

On the other hand, for a nonparametric Bootstrap, the procedure would be to draw new samples from the empirical distribution, i.e. simply draw samples with replacement from the initial sample X_1, \dots, X_n and calculate the new estimation. In this case, $P(\hat{\theta}^* = \hat{\theta}) \approx 1 - (1 - \frac{1}{n})^n$, since the event $\hat{\theta}^* = \hat{\theta}$ could be interpreted as the probability of $\hat{\theta}^*$ appearing at least once within the j -th resample for the nonparametric Bootstrap.

Now, in order to calculate the true probability of the event $\hat{\theta}^* = \hat{\theta}$, take the limit as $n \rightarrow \infty$. It is well known that $(1 - \frac{1}{n})^n \rightarrow e^{-1}$ as $n \rightarrow \infty$. Then, $P(\hat{\theta}^* = \hat{\theta}) = 1 - e^{-1} \approx 0.632$.

References

- [1] Robert Andersen. *Modern Methods for Robust Regression*. 152. Sage, 2008.
- [2] Spearman CE. “General Intelligence, Objectively Determined and Measured”. In: *American Journal of Psychology* 15 (1904), pp. 201–293.
- [3] Michael Falk. “On MAD and Comedians”. In: *Annals of the Institute of Statistical Mathematics* 49.4 (1997), pp. 615–644.
- [4] Maurice G Kendall. “A New Measure of Rank Correlation”. In: *Biometrika* 30.1/2 (1938), pp. 81–93.
- [5] R.J. Kulperger. “Parametric and Nonparametric Bootstrap”. In: *Handouts for Introduction to the Theory of Statistics* (2019). URL: <http://fisher.stats.uwo.ca/faculty/kulperger/SS3858/Handouts/Bootstrap.pdf>.
- [6] John Nachbar. “Concave and Convex Functions”. In: *Lecture Notes for Economics* 4111 (2018).

- [7] Joseph Lee Rodgers. “The Bootstrap, the Jackknife, and the Randomization Test: A Sampling Taxonomy”. In: *Multivariate Behavioral Research* 34.4 (1999), pp. 441–456.
- [8] Yuliana Susanti, Hasih Pratiwi, et al. “M Estimation, S Estimation, and MM Estimation in Robust Regression”. In: *International Journal of Pure and Applied Mathematics* 91.3 (2014), pp. 349–360.
- [9] Aad W Van der Vaart. *Asymptotic Statistics*. Vol. 3. Cambridge university press, 2000.
- [10] Larry Wasserman. *All of Nonparametric Statistics*. Springer Science & Business Media, 2006.